

As fast as it gets?

Challenging  (cpp) with 

Philipp Adämmer Philipp Wittenberg

St  alsund

11/2021



What is Julia?

- ▶ Compared to R, Julia¹ is a new programming language whose stable version has been released in 2018
- ▶ Julia is dynamically typed but also allows static type annotations
- ▶ Julia is fast and solves the two language problem
- ▶ The package environment is not as mature as R's but it is easy to call R or other languages
- ▶ Still suffers from compilation latency ("first time to plot")

¹(Bezanson et al. 2017; *SIAM review*)

Simple task

- ▶ Draw from $X \sim \text{Bin}(n = 100, p = 0.1)$
- ▶ Repeat the experiment $N = 10^7$ times and compute coverage rates for the following 95% confidence intervals:
 - Wald
 - Wilson
 - Clopper-Pearson
- ▶ Compare the performance between R(cpp) and Julia
- ▶ Start with a naive loop implementation

Use package to compute CIs

```
1  function f1(n, p, N,  $\alpha$ )
2
3  # Pre-allocate
4  waldci    = zeros(N, 2)
5  wilsonci  = zeros(N, 2)
6  cpci      = zeros(N, 2)
7
8  # Loop
9  Threads.@threads for jj = 1:N
10
11  # Compute boolean vector
12  randnrs = rand(n) .<= p
13
14  # Compute CIs
15  waldtemp    = confint(BinomialTest(randnrs, p), level =
16                        (1 -  $\alpha$ ), method = :wald)
17
18  wilsontemp  = confint(BinomialTest(randnrs, p), level =
19                        (1 -  $\alpha$ ), method = :wilson)
20
21  cptemp      = confint(BinomialTest(randnrs, p), level =
22                        (1 -  $\alpha$ ), method = :clopper_pearson)
23
24
25
```

Use package to compute CIs

```
26                                     .
27                                     .
28                                     .
29
30   # Fill matrices with CIs
31   waldci[jj, 1]    = waldtemp[1]
32   waldci[jj, 2]    = waldtemp[2]
33
34   wilsonci[jj, 1]  = wilsontemp[1]
35   wilsonci[jj, 2]  = wilsontemp[2]
36
37   cpci[jj, 1]      = cptemp[1]
38   cpci[jj, 2]      = cptemp[2]
39
40   end
41
42   # Compute coverage rates
43   [ sum(waldci[:, 1]    .<= p .<= waldci[:, 2])/N
44     sum(wilsonci[:, 1]  .<= p .<= wilsonci[:, 2])/N
45     sum(cpci[:, 1]      .<= p .<= cpci[:, 2])/N ]
46
47   end
```

Replace for loop

```
1  function f2(n, p, N,  $\alpha$ )
2
3
4  # Compute boolean Matrix
5      randnrs = rand(n, N) .<= p
6
7
8  # Compute confidence intervals
9      waldci = ThreadsX.map(x -> confint(BinomialTest(x, p),
10      level=(1 -  $\alpha$ ), method=:wald), eachcol(randnrs))
11
12      wilsonci = ThreadsX.map(x -> confint(BinomialTest(x, p),
13      level=(1 -  $\alpha$ ), method=:wilson), eachcol(randnrs))
14
15      cpci = ThreadsX.map(x -> confint(BinomialTest(x, p),
16      level=(1 -  $\alpha$ ), method=:clopper_pearson), eachcol(randnrs))
17
18
19  # Compute coverage rates
20      [ sum(first.(waldci) .<= p .<= last.(waldci))/N
21      sum(first.(wilsonci) .<= p .<= last.(wilsonci))/N
22      sum(first.(cpci) .<= p .<= last.(cpci))/N ]
23
24  end
```

What is Rcpp³?

- ▶ R-package providing functions for seamless integration of R and C++
- ▶ Brings speed and performance to address R's bottlenecks
- ▶ Currently used by 2442 CRAN packages
- ▶ Include C++ libraries and other APIs via `sourceCpp` or `cppFunction`
- ▶ Can be used for parallel computations via `RcppParallel`²

²(Allaire et al. 2021;)

³(Eddelbuettel and François. 2011; *J Stat Softw*)

Confidence intervals

- ▶ Wald interval

$$\hat{p} \pm z_{1-\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

- ▶ Wilson interval

$$\frac{1}{1 + \frac{z_{1-\alpha/2}^2}{n}} \left(\hat{p} + \frac{z_{1-\alpha/2}^2}{2n} \right) \pm \frac{z_{1-\alpha/2}}{1 + \frac{z_{1-\alpha/2}^2}{n}} \sqrt{\frac{\hat{p}(1-\hat{p})}{n} + \frac{z_{1-\alpha/2}^2}{4n^2}}$$

- ▶ Clopper–Pearson interval

$$\text{Beta}\left(\frac{\alpha}{2}; x, n - x + 1\right) < \theta < \text{Beta}\left(1 - \frac{\alpha}{2}; x + 1, n - x\right)$$

confidence intervals from scratch

```

1  #include <Rcpp.h>
2  using namespace Rcpp;
3
4  // [[Rcpp::export]]
5  NumericVector CIrcpp1(int n, double p, int N, double alpha) {
6
7      const double& a1 = alpha/2, a2 = 1-alpha/2;
8      const double& z = R::qnorm(a2, 0, 1, true, false), zz = z*z;
9
10     double x, c0, c1, b0, b1;
11     NumericMatrix CIV(N, 6);
12
13     for (int i = 0; i < N; i++) {
14         x = R::rbinom(n, p);
15
16         /* Wald CI */
17         c0 = x/n;
18         c1 = z*sqrt(c0*(1-c0)/n);
19         CIV(i, 0) = c0-c1;
20         CIV(i, 1) = c0+c1;
    
```

confidence intervals from scratch

```
21      /* Wilson CI */
22      b0 = (1/(1 + zz/n)) * (c0 + zz/(2*n));
23      b1 = z/(1 + zz/n) * sqrt( (c0*(1-c0)/n) + zz/(4*n*n) );
24      CIV(i, 2) = b0 - b1;
25      CIV(i, 3) = b0 + b1;
26
27      /* Clopper-Pearson CI */
28      CIV(i, 4) = R::qbeta(a1, x, n-x+1, true, false);
29      CIV(i, 5) = R::qbeta(a2, x+1, n-x, true, false);
30  }
31
32  /* compute coverage rates */
33  return NumericVector::create(
34      mean( (CIV(_, 0) <= p) & (p <= CIV(_, 1)) ),
35      mean( (CIV(_, 2) <= p) & (p <= CIV(_, 3)) ),
36      mean( (CIV(_, 4) <= p) & (p <= CIV(_, 5)) )
37  );
38 }
```

Rcpp and RcppParallel

```

1  #include <Rcpp.h>
2  #include <RcppParallel.h>
3  // [[Rcpp::depends(RcppParallel)]]
4  using namespace Rcpp;
5  using namespace RcppParallel;
6
7  struct CIworker : public Worker {
8
9      /* input to read from */
10     const RVector<double> X;
11     const double a1, a2, z;
12     const int n;
13
14     /* output matrix to write to */
15     RMatrix<double> mat;
16
17     /* initialize */
18     CIworker(const NumericVector X, double a1, double a2, double z, int n,
19             ↪ NumericMatrix mat) : X(X), a1(a1), a2(a2), z(z), n(n), mat(mat)
20
21     /* calculate intervals and write to output matrix */
22     .
23     .
24     .
25 };

```

Rcpp and RcppParallel

```

25  // [[Rcpp::export]]
26  NumericVector CIrcpp2(int n, double p, int N, double alpha) {
27
28      const double& a1 = alpha/2, a2 = 1-alpha/2;
29      const double& z = R::qnorm(a2, 0, 1, true, false);
30      /* draw vector of size N from Bin(n, p) */
31      const NumericVector& X(Rcpp::rbinom(N, n, p));
32
33      /* allocate the output matrix */
34      NumericMatrix CIV(N, 6);
35
36      /* create a worker */
37      CIworker f1(X, a1, a2, z, n, CIV);
38
39      /* call it with parallelFor */
40      parallelFor(0, N, f1);
41
42      /* return vector with coverage rates */
43      return NumericVector::create(
44          mean( (CIV(_, 0) <= p) & (p <= CIV(_, 1)) ),
45          mean( (CIV(_, 2) <= p) & (p <= CIV(_, 3)) ),
46          mean( (CIV(_, 4) <= p) & (p <= CIV(_, 5)) ));
47  }

```

Benchmarking

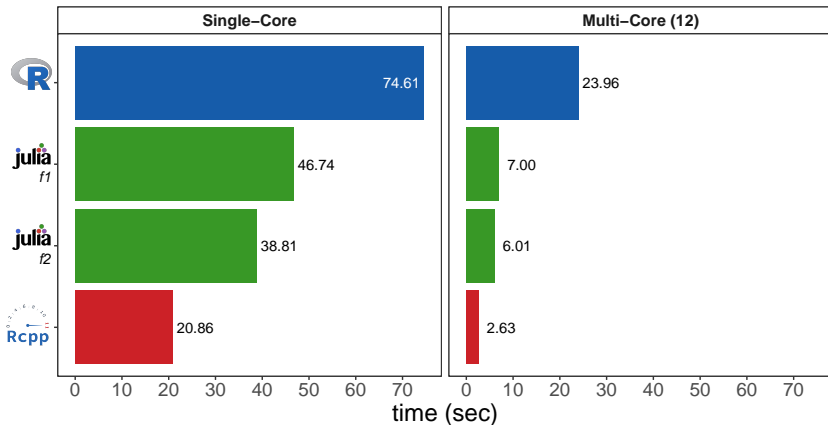
Rcpp and Julia

```
1  # Load packages
2  library(Rcpp)
3  library(RcppParallel)
4  library(microbenchmark)
5
6  # Source scripts
7  sourceCpp("singlecore.cpp")
8  sourceCpp("multicore.cpp")
9
10 # Set parameters
11 n    <- 100
12 p    <- .1
13 N    <- 1e7
14 alpha <- .05
15 setThreadOptions(numThreads = 12)
16
17 # Benchmark functions
18 microbenchmark(
19   CIrcpp1(n, p, N, alpha),
20   CIrcpp2(n, p, N, alpha),
21   times = 10)
```

```
1  # Load packages
2  using HypothesisTests
3  using ThreadsX
4  using BenchmarkTools
5
6  # Include scripts
7  include("Code1.jl")
8  include("Code2.jl")
9
10 # Set paramters
11 n = 100;
12 p = 0.1;
13 N = 107;
14 α = 0.05;
15
16 # Benchmark functions
17 @benchmark f1($n, $p, $N, $α) seconds = 70
18 @benchmark f2($n, $p, $N, $α) seconds = 70
```

Benchmark results

single and multi-core



Custom function to compute CIs

```
1  function ciallf(x::Int64, z::Float64, n::Int64,  
2                a1::Float64, a2::Float64)  
3  
4  
5      # Wald CI  
6      c0    = x/n  
7      c1    = z*sqrt(c0*(1 - c0)/n)  
8  
9  
10     # Wilson CI  
11     b0     = (1/(1 + z2/n))*(c0 + (z2)/(2*n))  
12     b1     = (z/(1 + (z2/n)))*sqrt((c0*(1 - c0)/n) + z2/(4*n2))  
13  
14  
15     # Clopper-Pearson CI  
16     cplow  = (x == 0 ? 0.0 : quantile(Beta(x,      n - x + 1), a1))  
17     cpup   = (x == n ? 1.0 : quantile(Beta(x + 1, n - x),  a2))  
18  
19  
20     # Return CIs  
21     [(c0 - c1) (c0 + c1) (b0 - b1) (b0 + b1) cplow cpup]  
22  
23 end
```

Use custom function and `Distributions.jl`⁴ to draw from `Bin(n, p)`

```
1  function f3(n::Int64, p::Float64, N::Int64, α::Float64)
2
3
4      # Compute quantile
5      z = quantile(Normal(), 1 - α/2)
6
7
8      # Draw # of successes
9      dist = Binomial(n, p)
10     randnrs = rand(dist, N)
11
12
13     # Compute CIs
14     ciall = ThreadsX.map(x -> ciallf(x, z, n, α/2, (1 - α/2)),
15                          randnrs)::Vector{Matrix{Float64}}
16
17
18     # Compute coverage rates
19     [ sum(getindex.(ciall, 1) .<= p .<= getindex.(ciall, 2))/N
20       sum(getindex.(ciall, 3) .<= p .<= getindex.(ciall, 4))/N
21       sum(getindex.(ciall, 5) .<= p .<= getindex.(ciall, 6))/N ]
22
23     end
```

⁴(Besançon et al. 2021; *J Stat Softw*)



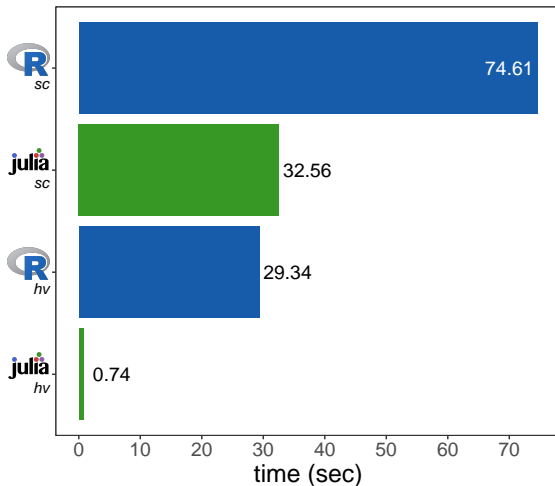
Version III

Use Static Arrays and draw CP before loop

```
1  function ciallf(x::Int64, z::Float64, n::Int64,  
2                cplow::Vector{Float64}, cpup::Vector{Float64},  
3                p::Float64)  
4  
5  # Wald CI  
6      c0    = x/n;  
7      c1    = z*sqrt(c0*(1 - c0)/n)  
8  
9  
10 # Wilson CI  
11     b0     = (1/(1 + z^2/n))*(c0 + (z^2)/(2*n))  
12     b1     = (z/(1 + (z^2/n)))*sqrt((c0*(1 - c0)/n) + z^2/(4*n^2))  
13  
14  
15 # Clopper-Pearson CIs already computed  
16         # --- #  
17  
18 # Return CIs  
19     SVector{3, Bool}((c0 - c1)      <= p <= (c0 + c1),  
20                      (b0 - b1)      <= p <= (b0 + b1),  
21                      cplow[x + 1]   <= p <= cpup[x + 1])  
22  
23 end
```

Benchmark results

Julia vs. R



sc: single core & full CI calculation;

hv: single core & half vectorized

unordered map

```

1  #include <Rcpp.h>
2  using namespace Rcpp;
3  // [[Rcpp::plugins(cpp11)]]
4
5  // [[Rcpp::export]]
6  NumericVector CIrcpp3(int n, double p, int N, double alpha) {
7
8      double c0, c1, b0, b1;
9      const double& a1 = alpha/2, a2 = 1-alpha/2;
10     const double& z = R::qnorm(a2, 0, 1, true, false), zz = z*z;
11     double WaL, WaU, WiL, WiU, L1, L2, L3;
12     std::unordered_map<int, double> CPL, CPU;
13
14     NumericVector X(Rcpp::rbinom(N, n, p));
15     NumericVector::iterator i;
16
17     /* unordered map for Clopper-Pearson CI values */
18     for (int j = 0; j <= n; j++) {
19         CPL[j] = R::qbeta(a1, j, n-j+1, true, false);
20         CPU[j] = R::qbeta(a2, j+1, n-j, true, false);
21     }

```

unordered map

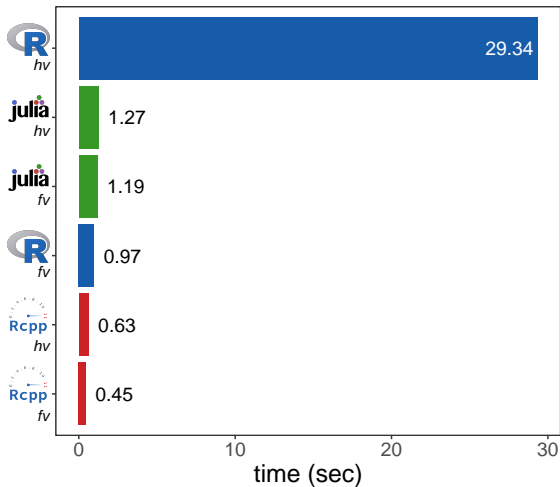
```

22  /* get intervals */
23  for (i = X.begin(); i != X.end(); ++i) {
24
25      /* Wald CI */
26      c0 = *i/n;
27      c1 = z*sqrt(c0*(1-c0)/n);
28      WaL = c0-c1;
29      WaU = c0+c1;
30
31      /* Wilson CI */
32      b0 = (1/(1 + zz/n)) * (c0 + zz/(2*n));
33      b1 = z/(1 + zz/n) * sqrt( (c0*(1-c0)/n) + zz/(4*n*n) );
34      WiL = b0-b1;
35      WiU = b0+b1;
36
37      if ( (WaL <= p) && (p <= WaU)) L1++;
38      if ( (WiL <= p) && (p <= WiU)) L2++;
39      if ((CPL[*i] <= p) && (p <= CPU[*i])) L3++;
40  }
41
42  /* compute coverage rates */
43  return NumericVector::create(L1/N, L2/N, L3/N);
44  }

```

Benchmarks

using on-board functions



hv: half vectorized


fv: full vectorized

Use RCall

```
1  function f5(n::Int64, p::Float64, N::Int64, α::Float64)
2
3
4
5      .
6      .
7      .
8
9
10
11
12     # Use R to draw from Bin(n, p)
13     randnrs::Vector{Int64} = (rcopy(Vector{Int64}, R"rbinom($N, $n, $p)"))
14
15
16
17     .
18     .
19     .
20
21
22
23  end
```

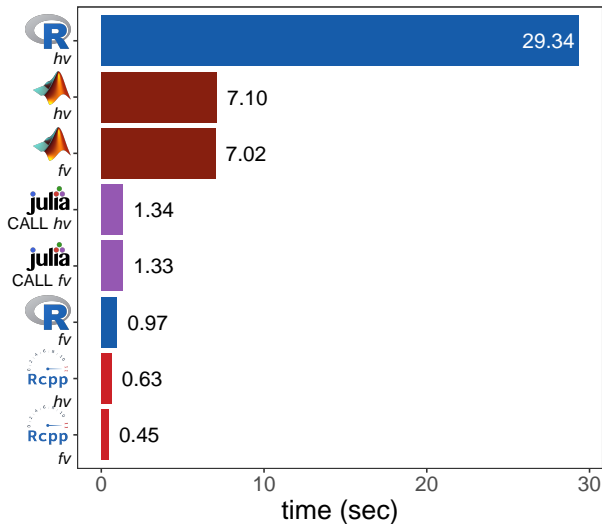
Benchmarks

using `rbinom()`



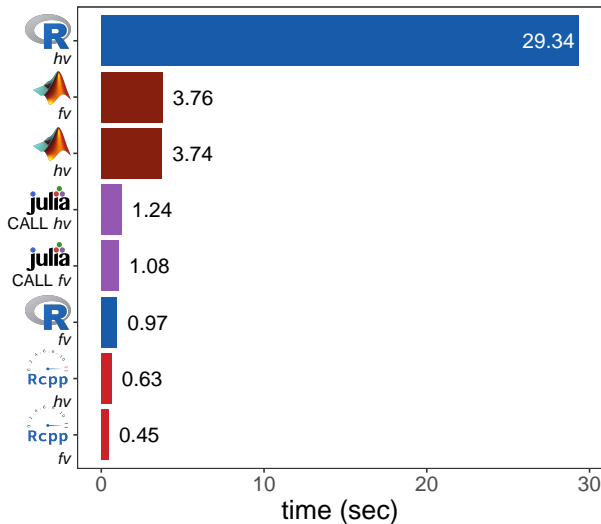
Additional benchmarks

using on-board functions



Additional benchmarks

using `rbinom()`



- ▶ Don't use loops in R for such a task, vectorize!
- ▶ Julia and Rcpp do not benefit much from vectorization
- ▶ Optimize your code before using multiple threads!
- ▶ Rcpp forces you to write type stable functions. Writing type stable functions in Julia can be tricky sometimes
- ▶ Using R within Julia is easier and faster than using Julia within R
- ▶ Using Rcpp is (slightly) faster for this task than using Julia(Call)

Thank you for your attention

References



Bezanson, J. et al. (2017). “Julia: A fresh approach to numerical computing”. *SIAM review* 59.1, pp. 65–98



Allaire, J. et al. (2021). *RcppParallel: Parallel Programming Tools for 'Rcpp'*. R package version 5.1.4. URL: <https://CRAN.R-project.org/package=RcppParallel>



Eddelbuettel, D. and R. François (2011). “Rcpp: Seamless R and C++ Integration”. *J Stat Softw* 40.8, pp. 1–18



Besaçon, M. et al. (2021). “Distributions.jl: Definition and Modeling of Probability Distributions in the JuliaStats Ecosystem”. *J Stat Softw* 98.16, pp. 1–30